

# Rekordbox Export Structure Analysis

James Elliott  
Deep Symmetry, LLC

March 3, 2019

## Abstract

The files written to external media by rekordbox for use in player hardware contain a wealth of information that can be used in place of queries to the remotedb server on the players, which is important because they can be obtained from the players' NFS servers, even if there are four players in use sharing the same media. Under those circumstances, remotedb queries are impossible. This article documents what has been learned so far about the files, and how to interpret them.

## Contents

<b>1</b>	<b>Database Exports</b>	<b>3</b>
1.1	File Header . . . . .	3
1.2	Table Pages . . . . .	6
1.3	Table Rows . . . . .	8
1.3.1	Album Rows . . . . .	8
1.3.2	Artist Rows . . . . .	9
1.3.3	Artwork Rows . . . . .	10
1.3.4	Color Rows . . . . .	10
1.3.5	Genre Rows . . . . .	10
1.3.6	Key Rows . . . . .	12
1.3.7	Label Rows . . . . .	12
1.3.8	Playlist Tree Rows . . . . .	12
1.3.9	Playlist Entry Rows . . . . .	13
1.3.10	Track Rows . . . . .	13
1.4	DeviceSQL Strings . . . . .	16
1.4.1	Long ASCII Strings . . . . .	17
1.4.2	Long UTF-16 Big-Endian Strings . . . . .	17
1.4.3	Short ASCII Strings . . . . .	17
<b>2</b>	<b>Analysis Files</b>	<b>18</b>

*CONTENTS*

2

**3   Crate Digger**

**18**

**List of Figures**

**18**

**List of Tables**

**19**

## 1 Database Exports

The starting point for finding track metadata from a player is the database export file, which can be found within rekordbox media at the path `PIONEER/rekordbox/export.pdb` (if you are using the Crate Digger FileFetcher to request this file, use that path as the `filePath` argument, and use a `mountPath` value of `/B/` if you want to read it from the SD slot, or `/C/` to obtain it from the USB slot).

The file is a relational database format designed to be efficiently used by very low power devices (there were deployments on 16 bit devices with 32K of RAM). Today you are most likely to encounter it within the Pioneer Professional DJ ecosystem, because it is the format that their rekordbox software uses to write USB and SD media which can be mounted in DJ controllers and used to play and mix music.

The file consists of a series of fixed size pages. The first page contains a file header which defines the page size and the locations of database tables of different types, by the index of their first page. The rest of the pages consist of the data pages for all of the tables identified in the header.

Each table is made up of a series of rows which may be spread across any number of pages. The pages start with a header describing the page and linking to the next page. The rest of the page is used as a heap: rows are scattered around it, and located using an index structure that builds backwards from the end of the page. Each row of a given type has a fixed size structure which links to any variable-sized strings by their offsets within the page.

As changes are made to the table, some records may become unused, and there may be gaps within the heap that are too small to be used by other data. There is a bit map in the row index that identifies which rows are actually present. Rows that are not present must be ignored: they do not contain valid (or even necessarily well-formed) data.

The majority of the work in reverse-engineering this format was performed by Henry Betts<sup>1</sup> and Fabian Lesniak<sup>2</sup>, to whom I am hugely grateful.

### 1.1 File Header

Unless otherwise stated, all multi-byte numbers in the file are stored in little-endian byte order. Field names used in the Figures match the IDs assigned to them in the Kaitai Struct specification<sup>3</sup>, unless that is

---

<sup>1</sup><https://github.com/henrybetts/Rekordbox-Decoding>

<sup>2</sup><https://github.com/flesniak/python-prodj-link>

<sup>3</sup>[https://github.com/Deep-Symmetry/crate-digger/blob/master/src/main/kaitai/rekordbox\\_pdb.ksy](https://github.com/Deep-Symmetry/crate-digger/blob/master/src/main/kaitai/rekordbox_pdb.ksy)




Figure 1: File Header




Figure 2: Table Pointer

too long to fit, in which case a subscripted abbreviation is used, and the text will mention the actual struct field name.

The first page begins with the file header, shown in Figure 1. The header starts with four zero bytes, followed by a four-byte integer, *len\_page* at byte 04, that establishes the size of each page (including this first one), in bytes. This is followed by another four-byte integer, *num\_tables* at byte 08, which reports the number of different tables that are present in the file. Each table will have a table pointer entry in the “Table pointers” section of the file header, described below, that identifies and locates the table.

The four-byte integer *next\_u* at byte 0c has an unknown purpose, but Mr. Lesniak named it *next\_unused\_page* and said “Not used as any *empty\_candidate*, points past the end of the file.” The four-byte integer *sequence*, at byte 14, was described “Always incremented by at least one, sometimes by two or three.” and I assume this means it reflects a version number that rekordbox updates when synchronizing to the exported media.

Finally, there is another series of four zero bytes, and then the header ends with the list of table pointers which begins at byte 1c. There are as many of these as specified by *num\_tables*, and each has the structure shown in Figure 2.

Each Table Pointer is a series of four four-byte integers. The first, *type*, identifies the type of table being defined. The known table types are shown in Table 1. The second value, at byte 04 of the Table Pointer, was called `empty_candidate` by Mr. Lesniak. It may link to a chain of empty pages if the database is ever garbage collected, but this is speculation on my part.

Type	Name	Row Content
0	tracks	Track metadata: title, artist, genre, artwork ID, playing time, etc., Section 1.3.10.
1	genres	Musical genres, for reference by tracks and searching, Section 1.3.5.
2	artists	Artists, for reference by tracks and searching, Section 1.3.2.
3	albums	Albums, for reference by tracks and searching, Section 1.3.1.
4	labels	Music labels, for reference by tracks and searching, Section 1.3.7.
5	keys	Musical keys, for reference by tracks, searching, and key matching, Section 1.3.6.
6	colors	Color labels, for reference by tracks and searching, Section 1.3.4.
7	playlist_tree	Describes the hierarchical tree structure of available playlists and folders grouping them, Section 1.3.8.
8	playlist_entries	Links tracks to playlists, in the right order, Section 1.3.9.
13	artwork	File paths of album artwork images, Section 1.3.3.
16	columns	Details not yet confirmed.
19	history	Records the tracks played during performances.

Table 1: Table Types

Other than the *type*, the two important values are *first\_page* at byte 08 and *last\_page* at byte 0c. These tell us how to find the table. They are page indices, where the page containing the file header has index 0, the page with index 1 begins at byte *len\_page*, and so on. In other words, the first page of the table identified by the current Table Pointer can be found within the file starting at the byte  $\text{len\_page} \times \text{first\_page}$ .

The table is a linked list of pages: each page contains the index of the next page after it. However, you need to keep track of the

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f							
00	00 00 00 00				page_index				type				next_page										
10	unknown <sub>1</sub>				unknown <sub>2</sub>				n <sub>rs</sub>	u <sub>3</sub>	u <sub>4</sub>	pf	free <sub>s</sub>	used <sub>s</sub>									
20	u <sub>5</sub>	num <sub>rl</sub>		u <sub>6</sub>		u <sub>7</sub>		heap															
row_groups																							
								...															
								ofs <sub>14</sub>		ofs <sub>13</sub>		ofs <sub>12</sub>		ofs <sub>11</sub>		ofs <sub>17</sub>		ofs <sub>16</sub>		row <sub>pf1</sub>		ofs <sub>15</sub>	
								ofs <sub>10</sub>		ofs <sub>9</sub>		ofs <sub>8</sub>		ofs <sub>7</sub>									
								ofs <sub>6</sub>		ofs <sub>5</sub>		ofs <sub>4</sub>		ofs <sub>3</sub>		ofs <sub>2</sub>		ofs <sub>1</sub>		ofs <sub>0</sub>		row <sub>pf0</sub>	

Figure 3: Table Page

*last\_page* value for the table, because it tells you not to try to follow the next page link once you reach the page with that index. (If you do keep going, you will start reading pages of some different table.) The structure of the table pages themselves are described in the next section.

As far as we know, the remainder of the first page after the table pointers is unused.

## 1.2 Table Pages

The table header is followed by the table pages themselves. These each have the size specified by *len\_page* in Figure 1, and the structure shown in Figure 3.

Data pages all seem to have the header structure described here, but not all of them actually store data. Some of them are “strange” and we have not yet figured out why. The discussion below describes how to recognize a strange page, and avoid trying to read it as a data page.

The first four bytes of a table page always seem to be zero. This is followed by a four-byte value *page\_index* which identifies the index of this page within the list of table pages (the header has index 0, the first actual data page the index 1, and so on). This value seems to be

redundant, because it can be calculated by dividing the offset of the start of the page by *len\_page*, but perhaps it serves as a sanity check.

This is followed by another four-byte value, *type*, which identifies the type of the page, using the values shown in Table 1. This again seems redundant because the table header which was followed to reach this page also identified the table type, but perhaps it is another sanity check, or an alternate way to tell, when following page links, that you have reached the end of the table you are interested in. Speaking of which, the next four-byte value, *next\_page*, is that link: it identifies the index at which the next page of this table can be found, as long as we have not already reached the final page of the table, as described in Section 1.1.

The exact meaning of *unknown<sub>1</sub>* is unclear. Mr. Flesniak said “sequence number (0→1: 8→13, 1→2: 22, 2→3: 27)” but I don’t know how to interpret that. Even less is known about *unknown<sub>2</sub>*. But *num\_rows\_small* at byte 18 within the page (abbreviated *n<sub>rs</sub>* in Figure 3) holds the number of rows that are present in the page, unless *num\_rows\_large* (below) holds a value that is larger (but not equal to 1fff). This seems like a strange mechanism for dealing with the fact that some tables (like playlist entries) have a lot of very small rows, too many to count with a single byte. But then why not just always use *num\_rows\_large*?

The purpose of the next two bytes are also unclear. Of *u<sub>3</sub>* Mr. Flesniak said “a bitmask (first track: 32)”, and he described *u<sub>4</sub>* as often 0, sometimes larger, especially for pages with a high number of rows (e.g. 12 for 101 rows).

Byte 1b is called *page\_flags* (abbreviated *p<sub>f</sub>* in Figure 3). According to Mr. Flesniak, “strange” (non-data) pages will have the value 44 or 64, and other pages have had the values 24 or 34. Crate Digger considers a page to be a data page if *page\_flags* & 40 = 0.

Bytes 1c–1d are called *free\_size* (abbreviated *free<sub>s</sub>* in Figure 3), and store the amount of unused space in the page heap (excluding the row index which is built backwards from the end of the page); *used\_size* at bytes 1c–1d (abbreviated *used<sub>s</sub>*) stores the number of bytes that are in use in the page heap.

Bytes 20–21, *u<sub>5</sub>*, are of unclear purpose. Mr. Flesniak labeled them “(0→1: 2).”

Bytes 20–21, *num\_rows\_large* (abbreviated *num<sub>rl</sub>* in Figure 3) hold the number of entries in the row index at the end of the page when that value is too large to fit into *num\_rows\_small* (as mentioned above), and that situation seems to be indicated when this value is larger than *num\_rows\_small*, but not equal to 1fff.

*u<sub>6</sub>* at bytes 24–25 seems to have the value 1004 for strange pages, and 0000 for data pages. And Mr. Flesniak describes *u<sub>7</sub>* at bytes 26–27 as “always 0 except 1 for history pages, num entries for strange pages?”

After these header fields comes the page heap. Rows are allocated within this heap starting at byte 28. Since rows can be different sizes, there needs to be a way to locate them. This takes the form of a row index, which is built from the end of the page backwards, in groups of up to sixteen row pointers along with a bitmask saying which of those rows are still part of the table (they might have been deleted).

The number of row index entries is determined by the value of either *num\_rows\_small* or *num\_rows\_large* as described above.

The bit mask for the first group of up to sixteen rows, labeled *row\_pf<sub>0</sub>* in Figure 3 (meaning “row presence flags group 0”), is found in the last two bytes of the page. The low order bit of this value will be set if row 0 is really present, the next bit if row 1 is really present, and so on. The two bytes before these flags, labeled *ofs<sub>0</sub>*, store the offset of the first row in the page. This offset is the number of bytes past the end of the page header at which the row itself can be found. So if row 0 begins at the very beginning of the heap, at byte 28 in the page, *ofs<sub>0</sub>* would have the value 0000.

As more rows are added to the page, space is allocated for them in the heap, and additional index entries are added at the end of the heap, growing backwards. Once there have been sixteen rows added, all of the bits in *row\_pf<sub>0</sub>* are accounted for, and when another row is added, before its offset entry *ofs<sub>16</sub>* can be added, another row bit-mask entry *row\_pf<sub>1</sub>* needs to be allocated. And so the row index grows backwards towards the rows that are being added forwards, and once they are too close for a new row to fit, the page is full, and another page gets allocated to the table.

### 1.3 Table Rows

The structure of the rows themselves is determined by the *type* of the table, using the values shown in Table 1.

#### 1.3.1 Album Rows

Album rows hold an album name and ID along with an artist association, with the structure shown in Figure 4. The unknown value at bytes 00-01 seems to usually have the values 80 00. It is followed by a two-byte value Mr. Flesniak called *index\_shift*, although I don’t know what that means, and another four bytes of unknown purpose. But at bytes 08-0b we finally find a value we have a use for: *artist\_id* holds the ID of an artist row associated with this track row. This is followed by *id*, the ID of this track row itself, at bytes 0c-0f. We assume that there are index tables somewhere that would let us locate the page and row index of a record given its table type and ID, but we have not yet found and figured them out.



	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
00	$u_1$		$i_{shift}$		$unknown_2$			$artist\_id$			$id$					
10	$unknown_3$				$u_4$	$o_n$										

Figure 4: Album Row

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
	0060		$i_{shift}$		$id$			$u_1$	$o_n$							

Figure 5: Artist Row with Nearby Name

This is followed by five more bytes with unknown meaning, and the final byte in the row,  $ofs\_name$  is a pointer to the track name (labeled  $o_n$  in Figure 4). To find the location of the name, add  $ofs\_name$  bytes to the address of the start of the track row itself. The name itself is encoded in a surprisingly baroque way, explained in Section 1.4

### 1.3.2 Artist Rows

Artist rows hold an Artist name and ID, with the structure shown in Figure 5 or Figure 6. The *subtype* value at bytes 00–01 determines which variant is used. If the artist name was allocated close enough to the row to be reached by a single byte offset, offset, *subtype* has the value 0060, and the row has the structure in Figure 5. If the name is too far away for that, *subtype* has the value 0064 and the row has the structure in Figure 6.

In either case, *subtype* is followed by the unexplained two-byte value found in many row types that Mr. Flesniak called *index\_shift*, and then by *id*, the ID of this artist row itself, at bytes 04–07, an unknown value, and *ofs\_name\_near* (labeled  $o_n$  in Figure 5), the one-byte name offset used only in the first variant.

If *subtype* is 0064, the value of *ofs\_name\_near* is ignored, and instead the two-byte value *ofs\_name\_far* (labeled  $o_{far}$  in Figure 6) is used.

Whichever name offset is used, it is a pointer to the artist name. To find the location of the name, add the value of the offset to the address of the start of the artist row itself. This gives the address of a DeviceSQL string holding the name, with the structure explained in Section 1.4

0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
0064	<i>i<sub>shift</sub></i>			<i>i<sub>d</sub></i>			<i>u<sub>1</sub></i>	<i>o<sub>n</sub></i>	<i>o<sub>far</sub></i>						

Figure 6: Artist Row with Far Name




Figure 7: Artwork Row

### 1.3.3 Artwork Rows

Artwork rows hold an *id* (which tracks refer to) and the path at which the corresponding album art image file can be found, with the structure shown in Figure 7. Note that in this case, the DeviceSQL string *path* is embedded directly into the row itself, rather than being located elsewhere in the heap through an offset. The structure of the string itself is still as described in Section 1.4.

### 1.3.4 Color Rows

Color rows hold a numeric color *id* (which controls the actual color displayed on the player interface) at bytes 05–06 and a text label or *name* starting at byte 08 which is a DeviceSQL string shown in the information panel for tracks that are assigned the color. The rows have the structure shown in Figure 8. There are several bytes in the row that are not yet known to have any meaning. The structure of *name* is described in Section 1.4.

### 1.3.5 Genre Rows

Genre rows hold a numeric genre *id* (which tracks can be assigned) at bytes 00–03 and a text *name* starting at byte 04 which is a DeviceSQL string. The rows have the structure shown in Figure 9. The structure of *name* is described in Section 1.4.




Figure 8: Color Row




Figure 9: Genre or Label Row




Figure 10: Key Row

### 1.3.6 Key Rows

Key rows represent musical keys. They hold a numeric *id* (which tracks can be assigned) at bytes 00–03 and a text *name* starting at byte 08 which is a DeviceSQL string. (There seems to be a second copy of the ID at bytes 04–07.) The rows have the structure shown in Figure 10. The structure of *name* is described in Section 1.4.

### 1.3.7 Label Rows

Label rows represent record labels. They hold a numeric *id* (which tracks can be assigned) at bytes 00–03 and a text *name* starting at byte 04 which is a DeviceSQL string. The rows have the structure shown in Figure 9. The structure of *name* is described in Section 1.4.

### 1.3.8 Playlist Tree Rows

Playlist tree rows are used to organize the hierarchical structure of the playlist menu. There is probably an index somewhere that makes it possible to find the right rows directly when loading a playlist, but we have not yet figured out how indices work in DeviceSQL databases, so Crate Digger simply reads all the rows and builds its own in-memory index of the tree.

Playlist tree rows can either represent a playlist “folder” which contains other folders and playlists, or a regular playlist which holds only tracks. The rows are identified by an *id* at bytes 0c–0f, and also contain a *parent\_id* at bytes 00–03 which is how the hierarchical structure is represented: the contents of a folder are the other rows in this table whose *parent\_id* folder is equal to the *id* of the folder.

Similarly, the tracks that make up a regular playlist are the Playlist Entry Rows (described in Section 1.3.9) whose *playlist\_id* is equal to this row’s *id*.




Figure 11: Playlist Tree Row

Each playlist tree row also has a text *name* starting at byte 14 which is a DeviceSQL string displayed when navigating the hierarchy, a *sort\_order* indicator at bytes 08–0b (this may be the same value used to select sort orders when requesting menus using the `dbserver` protocol, shown in Table 7 of the dysentery Protocol Analysis paper, but this has not yet been confirmed), and a value that specifies whether the row defines a folder or a playlist. In the Kaitai Struct, this value is called *raw\_is\_folder*, is found at bytes 10–13, and has a non-zero value for folders. For convenience, the struct also defines a derived value, *is\_folder*, which is a boolean.

The rows have the structure shown in Figure 11. The structure of *name* is described in Section 1.4.

### 1.3.9 Playlist Entry Rows

Playlist entry rows list the tracks that belong to a particular playlist, and also establish the order in which they should be played. They have a very simple structure, shown in Figure 12, containing only three values. The *entry\_index* at bytes 00–03 specifies the position within the playlist at which this entry belongs. The *track\_id* at bytes 04–07 identifies the track to be played at this position in the playlist, by corresponding to the *id* of a row in the Track table described in Section 1.3.10, and the *playlist\_id* at bytes 08–0b identifies the playlist to which it belongs, by corresponding to the *id* of a row in the Playlist Tree described in Section 1.3.8.

### 1.3.10 Track Rows

Track rows describe audio tracks that can be played from the media export, and provide many details about the music including links to

0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
<i>entry_index</i>				<i>track_id</i>				<i>playlist_id</i>							

Figure 12: Playlist Entry Row

	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
00	$u_1$		$i_{shift}$		$bitmask$			$sample\_rate$				$composer\_id$				
10	$file\_size$				$u_2$			$u_3$		$u_4$		$artwork\_id$				
20	$key\_id$				$orig\_artist\_id$			$label\_id$				$remixer\_id$				
30	$bitrate$				$track\_number$			$tempo$				$genre\_id$				
40	$album\_id$				$artist\_id$			$id$				$disc_n$		$play_c$		
50	$year$		$s_{depth}$		$dur$		$u_5$		$c_{id}$	$r$	$u_6$		$u_7$		$ofs_0$	
60	$ofs_1$		$ofs_2$		$ofs_3$		$ofs_4$		$ofs_5$		$ofs_6$		$ofs_7$		$ofs_8$	
70	$ofs_9$		$ofs_{10}$		$ofs_{11}$		$ofs_{12}$		$ofs_{13}$		$ofs_{14}$		$ofs_{15}$		$ofs_{16}$	
70	$ofs_{17}$		$ofs_{18}$		$ofs_{19}$		$ofs_{20}$									

Figure 13: Track Row

other tables like artists, albums, keys, and others. They have the structure shown in Figure 13.

The first two bytes, labeled  $u_1$ , have an unknown purpose; they usually are 24 followed by 00. They are followed by the unexplained two-byte value found in many row types that Mr. Flesniak called *index\_shift*, and a four-byte value he called *bitmask*, although we do not know what the bits mean. The value *sample\_rate* at bytes 08–0b is the first one we have a solid understanding of: it holds the playback sample rate of the audio file, in samples per second.

Bytes 0c–0f hold the value *composer\_id* which identifies the composer of the track, if known, as a non-zero *id* value of an Artist row, as discussed in Section 1.3.2. The size of the audio file, in bytes, is found in *file\_size* at bytes 10–13. This is followed by an unknown four-byte value,  $u_2$ , which may be another ID, and two unknown two-byte values,  $u_3$  (about which Mr. Flesniak says “always 19048?”) and  $u_4$  (“always 30967?”).

If there is cover art for the track, there will be a non-zero value in *artwork\_id* (bytes 1c–1f), identifying the *id* of an Artwork row, as

discussed in section 1.3.3.

If a dominant musical key was identified for the track there will be a non-zero value in *key\_id* (bytes 20–23), which represents the *id* of a Key row, as discussed in Section 1.3.6. If the track is known to be a remake, the non-zero Artist row *id* (Section 1.3.2) of the original performer will be found in *original\_artist\_id* at bytes 24–27. If there is a known record label for the track, the non-zero value in *label\_id* will link to the *id* of a Label row *id* as described in Section 1.3.7. Similarly, if there is a known remixer, there will be a non-zero value in *remixer\_id* (bytes 2c–2f) linking to the *id* of an Artist row (Section 1.3.2).

The field *bitrate* at bytes 30–33 stores the playback bit rate of the track, and *track\_number* at bytes 34–37 holds the position of the track within its album. *tempo* at bytes 38–3b holds the playback tempo of the start of the track in beats per minute, multiplied by 100 (in order to support a precision of  $\frac{1}{100}$  BPM). If there is a known genre for the track, there will be a non-zero value in *genre\_id* at bytes 3c–3f, representing the *id* of a Genre row as discussed in Section 1.3.5.

If the track is part of an album, there will be a non-zero value in *album\_id* at bytes 40–43, and this will be the *id* of an Album row as discussed in Section 1.3.1. The Artist row *id* (Section 1.3.2) of the primary performer associated with the track is found in *artist\_id* at bytes 44–47. And the *id* of the track itself is found in *id* at bytes 48–4b. If the album is known to consist of multiple discs, the disc number on which this track is found will be in *disc\_number* at bytes 4c–4d. And the number of times the track has been played is found in *play\_count* (bytes 4e–4f).

The year in which the track was recorded, if known, is in *year* at bytes 50–51. The sample depth of the track audio file (bits per sample) is in *sample\_depth* at bytes 52–53. The playback time of the track (in seconds, at normal speed) is in *duration* at bytes 54–55. The purpose of the next two bytes, labeled  $u_5$ , is unknown, they seem to always hold the value 41.

Byte 58, *color\_id* (labeled  $c_{id}$  in Figure 13), holds the color assigned to the track in rekordbox, as the *id* of a Color row (described in Section 1.3.4), or zero if no color has been assigned. Byte 59, *rating* (labeled  $r$  in Figure 13) holds the rating (0 to 5 stars) assigned the track. The next two bytes, labeled  $u_6$ , have an unknown purpose, and seem to always have the value 1. The two bytes after them, labeled  $u_7$ , are also unknown; Mr. Flesniak said “alternating 2 and 3”.

The rest of the track row is an array of 21 two-byte offsets that point to DeviceSQL strings. To find the start of the string, add the address of the start of the track row to the offset. The purpose of each string is described in Table 2, and the structure of the strings themselves is explained in Section 1.4.

For convenience, the strings can be accessed as Kaitai Struct in-

stance values with the names shown in the table.

Index	Name	Content
0	<i>unknown_string_1</i>	Unknown, so far always empty.
1	<i>texter</i>	Unknown, named by @flesniak.
2	<i>unknown_string_2</i>	Unknown, “thought track number, wrong”.
3	<i>unknown_string_3</i>	Unknown, “Strange strings.” <sup>4</sup> .
4	<i>unknown_string_4</i>	“Strange strings” (as above).
5	<i>message</i>	Unknown, named by @flesniak.
6	<i>kuvo_public</i>	Empty or "ON". <sup>5</sup>
7	<i>autoload_hotcues</i>	Empty or "ON". <sup>6</sup>
8	<i>unknown_string_5</i>	Unknown.
9	<i>unknown_string_6</i>	Unknown, usually empty.
10	<i>date_added</i>	When track was added to collection.
11	<i>release_date</i>	When track was released.
12	<i>mix_name</i>	Name of the track remix.
13	<i>unknown_string_7</i>	Unknown, usually empty.
14	<i>analyze_path</i>	File path of track analysis, see Section 2.
15	<i>analyze_date</i>	When track analysis was performed.
16	<i>comment</i>	Track comment assigned by DJ.
17	<i>title</i>	Track title.
18	<i>unknown_string_8</i>	Unknown, usually empty.
19	<i>filename</i>	Name of track audio file.
20	<i>file_path</i>	File path of track audio.

Table 2: Track Offset Strings

## 1.4 DeviceSQL Strings

Many row types store string values, sometimes by directly embedding them, but more often by storing an offset to a location elsewhere in the heap. In either case the string itself uses the strange structure described in this section. Strings can be stored in a variety of formats. The first byte of the structure, labeled *length\_and\_kind* in the parsed Kaitai Struct, identifies the encoding type and, when the value is odd, also the length (for short ASCII strings), as detailed in Section 1.4.3.

<sup>4</sup>Often zero length, sometimes low binary values, ASCII 01 or 02 as content.

<sup>5</sup>Apparently used rather than a simple bit flag to control whether the track information is visible on Kuvo.

<sup>6</sup>Apparently used rather than a simple bit flag to control whether hot cues are automatically loaded for the track.






Figure 14: Long ASCII DeviceSQL String




Figure 15: Long UTF-16-BE DeviceSQL String

#### 1.4.1 Long ASCII Strings

If *length\_and\_kind* has the value 40, it is followed by a two-byte *length* field, and then followed by that many bytes of ASCII-encoded string data, as shown in Figure 14.

#### 1.4.2 Long UTF-16 Big-Endian Strings

If *length\_and\_kind* has the value 90, it is followed by a two-byte *length* field, and then followed by that many bytes of UTF-16 big-endian encoded string data, as shown in Figure 15.

#### 1.4.3 Short ASCII Strings

If *length\_and\_kind* has an odd value it is a *mangled\_length*, labeled  $m_l$  in Figure 16. This means we are dealing with a short ASCII DeviceSQL string. To find the length of the string data (which immediately follows this byte), subtract 1 from *mangled\_length*, divide it by 2, and subtract




Figure 16: Short ASCII DeviceSQL String

1 again.

2 Analysis Files

To be written...

3 Crate Digger

You can find a Java library that can parse the structures described in this research, and that can retrieve them from players' NFS servers, at: <https://github.com/deep-symmetry/crate-digger>

The project also contains Kaitai Struct specifications for the file structures, which were used to automatically generate Java classes to parse them, and which can be used to generate equivalent code for a variety of other programming languages.

There are also ONC RPC specification files which were similarly used to generate Java classes to communicate with the NFSv2 servers in the players, and which can likely be used to generate structures for other languages as well.

List of Figures

1	File Header . . . . .	4
2	Table Pointer . . . . .	4
3	Table Page . . . . .	6
4	Album Row . . . . .	9
5	Artist Row with Nearby Name . . . . .	9
6	Artist Row with Far Name . . . . .	10
7	Artwork Row . . . . .	10

8	Color Row . . . . .	11
9	Genre or Label Row . . . . .	11
10	Key Row . . . . .	12
11	Playlist Tree Row . . . . .	13
12	Playlist Entry Row . . . . .	14
13	Track Row . . . . .	14
14	Long ASCII DeviceSQL String . . . . .	17
15	Long UTF-16-BE DeviceSQL String . . . . .	17
16	Short ASCII DeviceSQL String . . . . .	18

## List of Tables

1	Table Types . . . . .	5
2	Track Offset Strings . . . . .	16



<http://deepsymmetry.org>